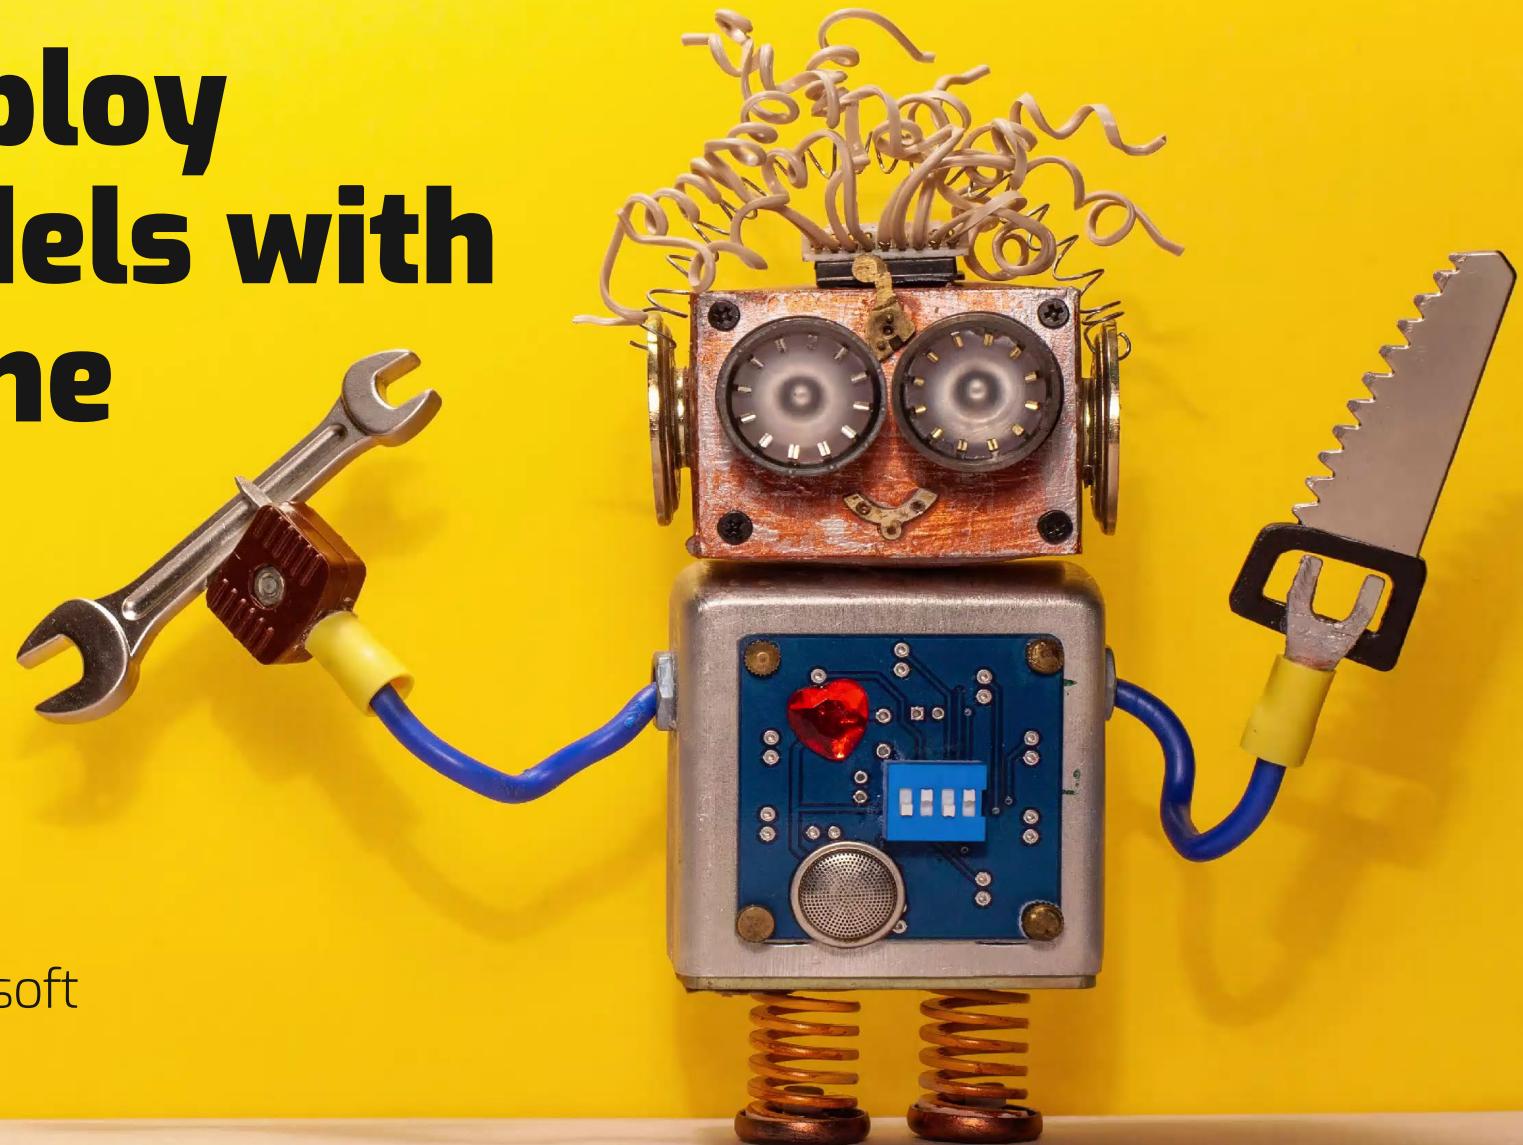


Build and deploy PyTorch models with Azure Machine Learning



Henk Boelman
Cloud Advocate @ Microsoft



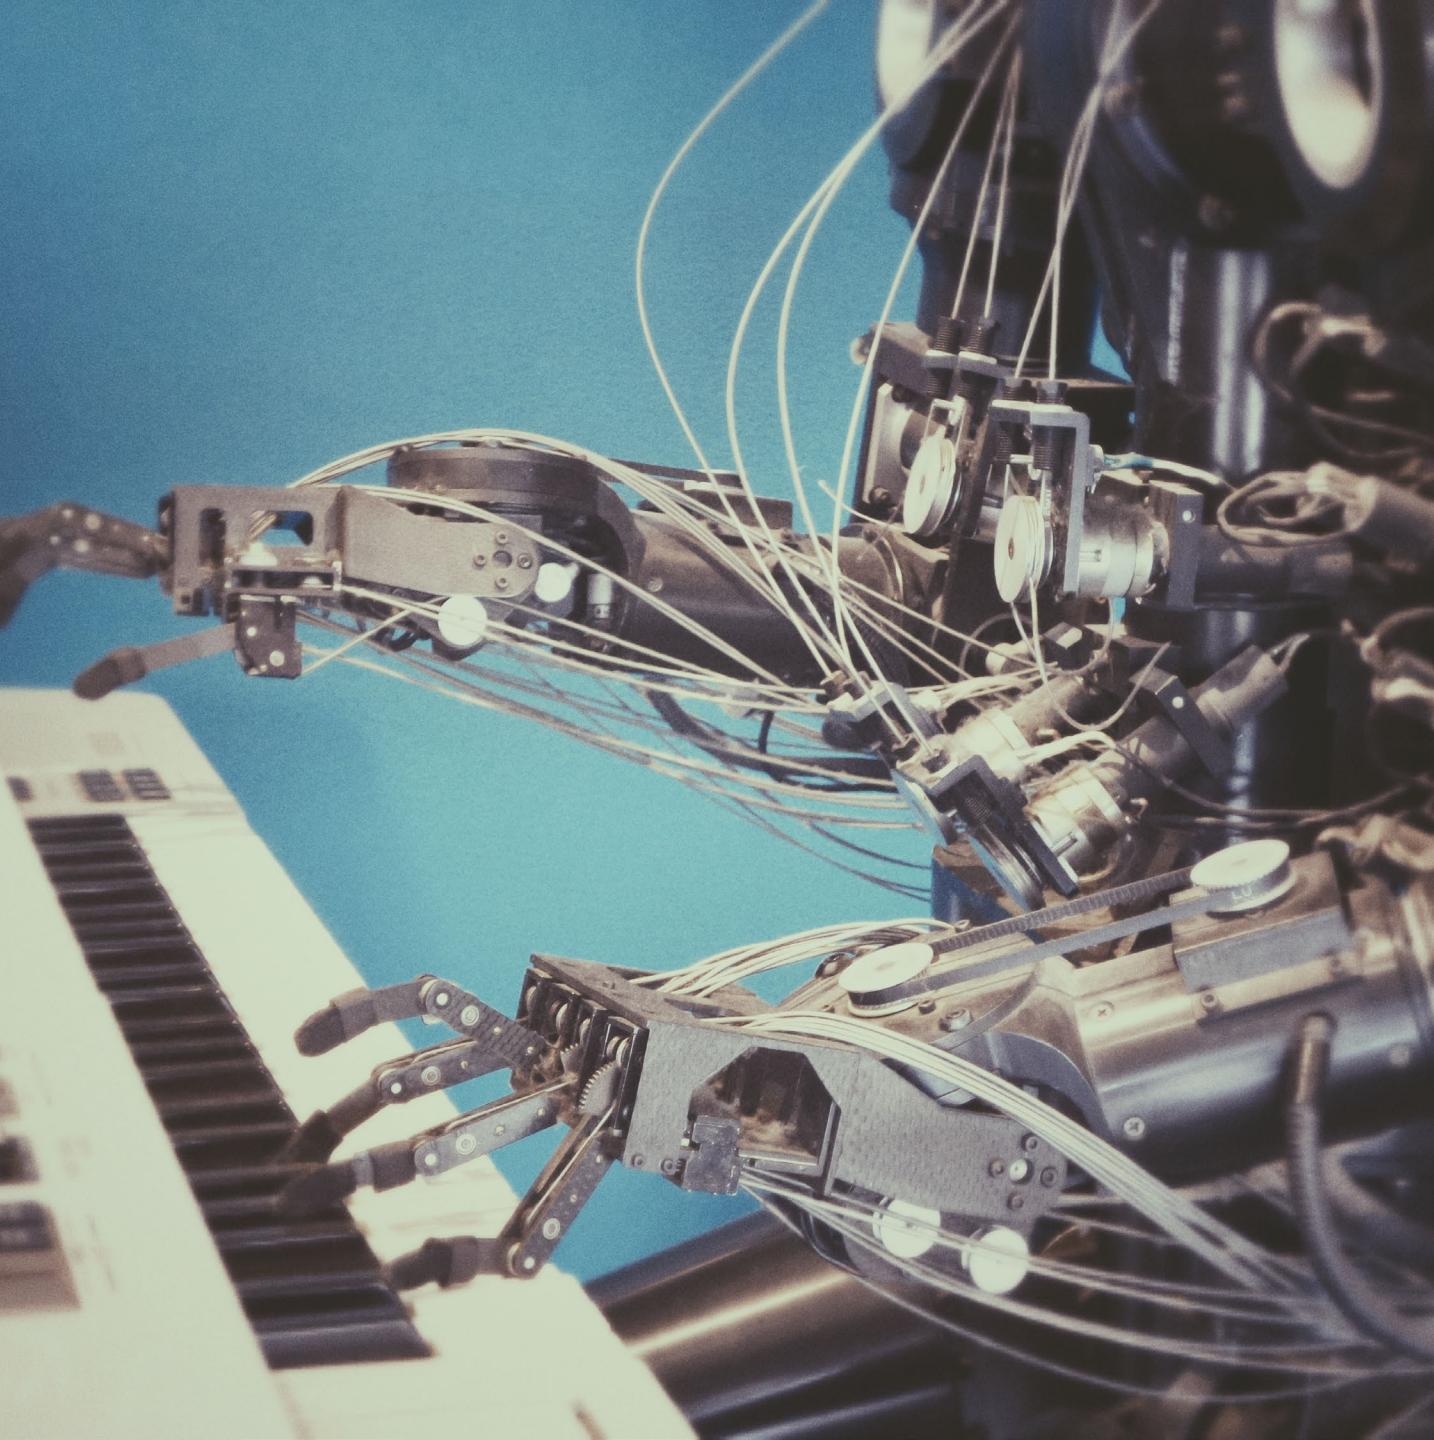
Agenda

Machine Learning on Azure
Simpson classification using Azure Machine Learning



Machine Learning

Ability to learn without being explicitly programmed.



Programming



Machine Learning



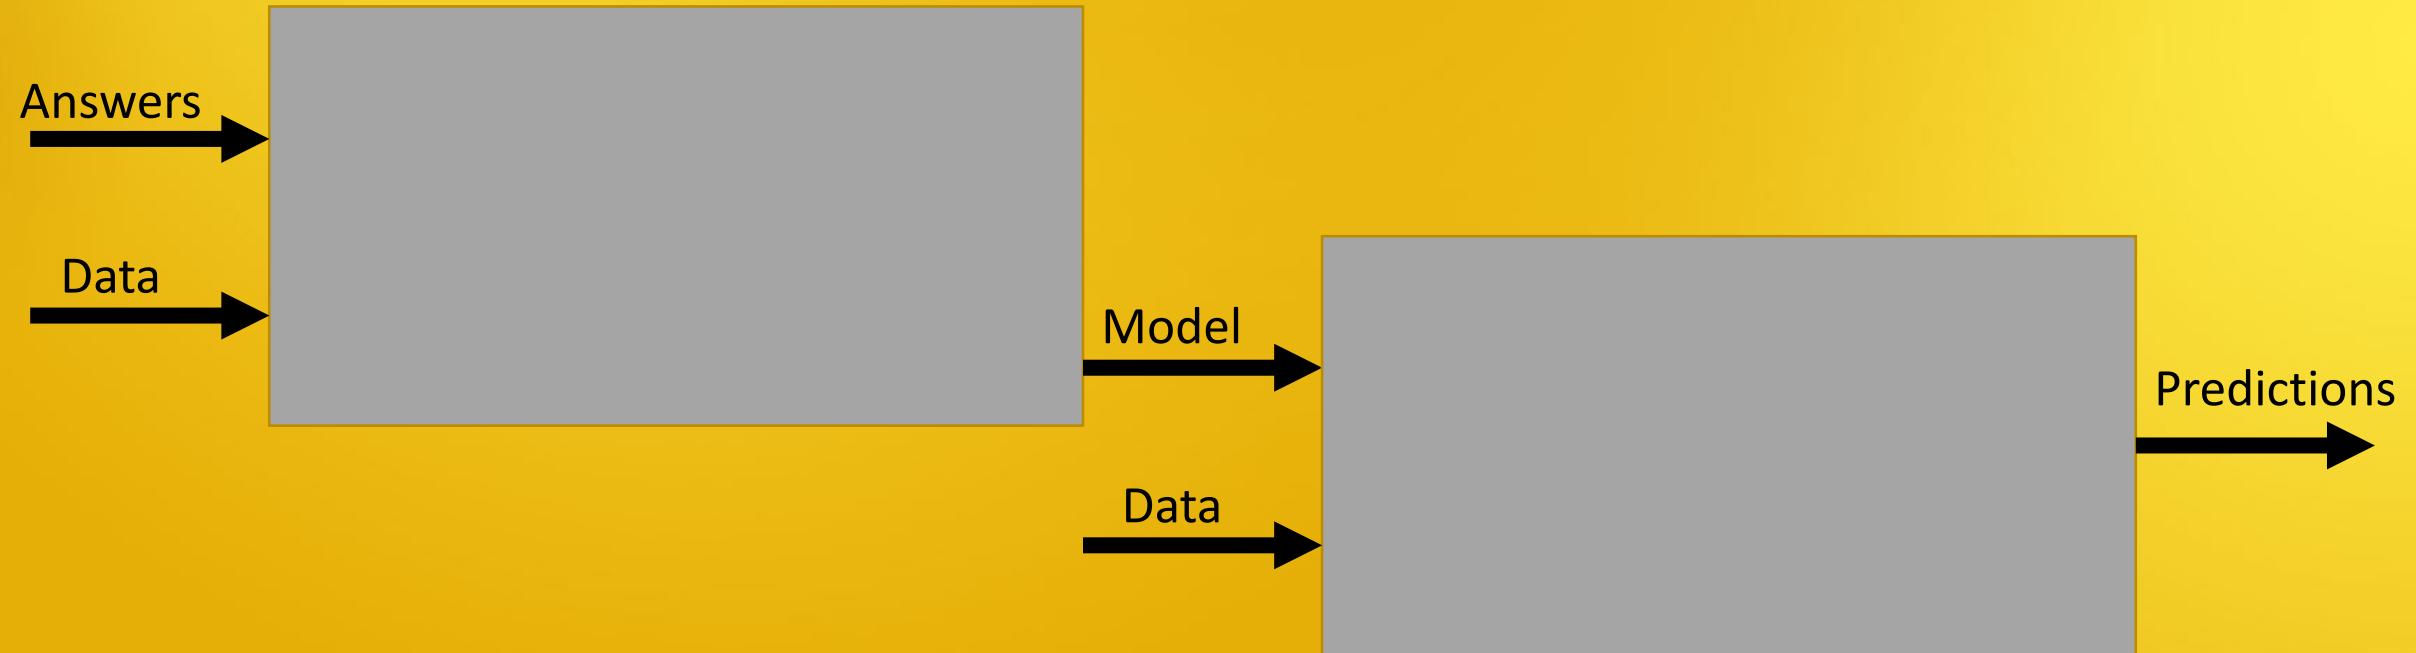
Machine Learning



Machine Learning



Machine Learning



Machine Learning on Azure

Sophisticated pretrained models

To simplify solution development

Cognitive Services



Vision



Speech



Language



Azure Search

Popular frameworks

To build advanced deep learning solutions



Pytorch



TensorFlow



Keras



Onnx

Productive services

To empower data science and development teams



Azure Databricks



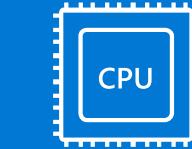
Azure Machine Learning



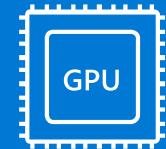
Machine Learning VMs

Powerful infrastructure

To accelerate deep learning



CPU



GPU



FPGA

Flexible deployment

To deploy and manage models on intelligent cloud and edge



On-premises



Cloud



Edge

Cognitive Services:

Pre-Trained models in the cloud and on the edge



Vision



Speech



Language



Knowledge

LEVEL UP!





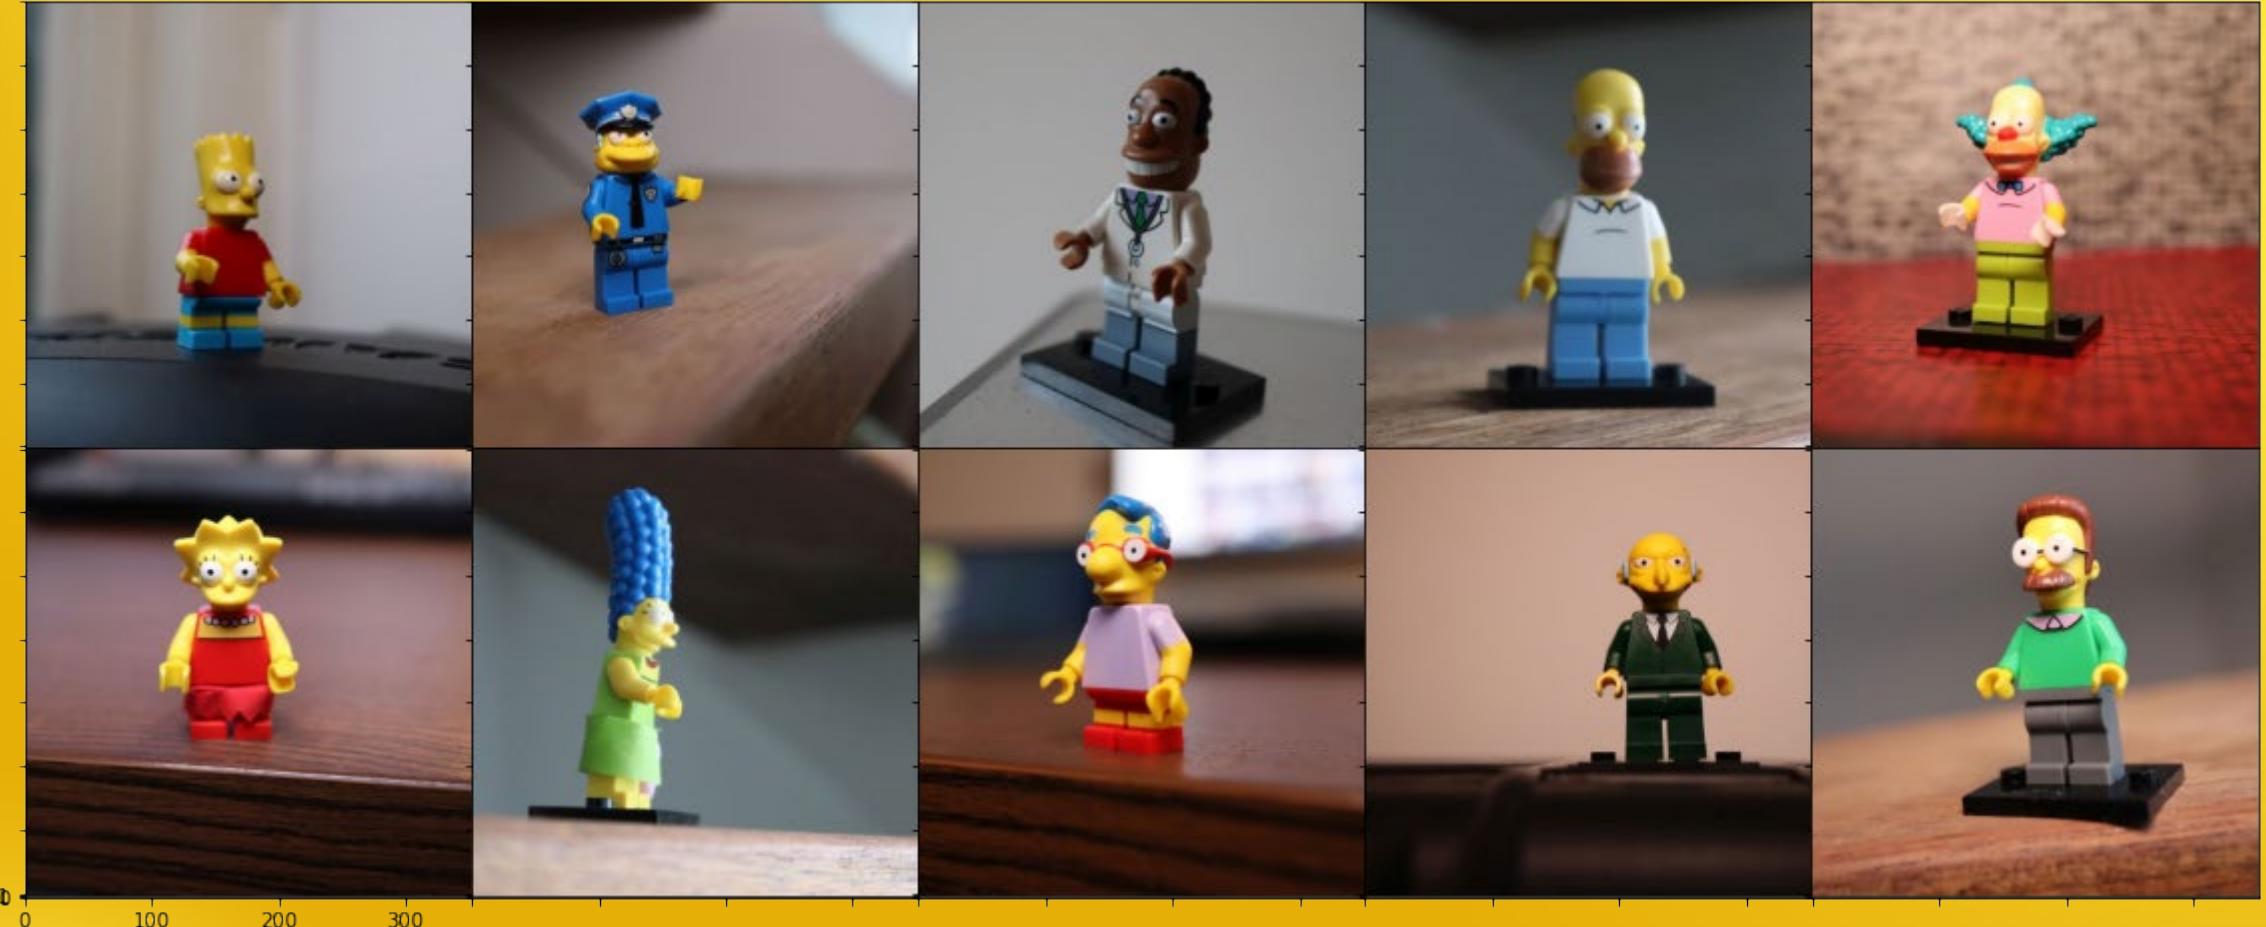
Azure Machine Learning studio

A fully-managed cloud service that enables you to easily build, deploy, and share predictive analytics solutions.



Is it Marge or Homer or .. or ..?

Simpson Lego dataset



<https://github.com/hnky/dataset-lego-figures>

Machine Learning on Azure

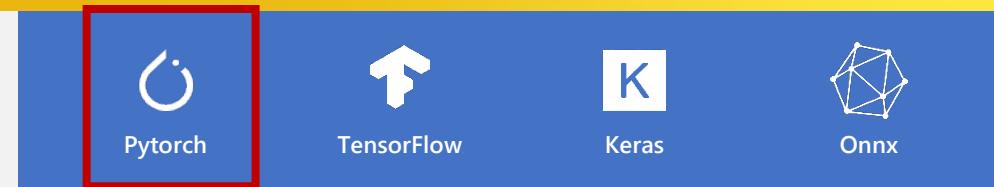
Sophisticated pretrained models

To simplify solution development



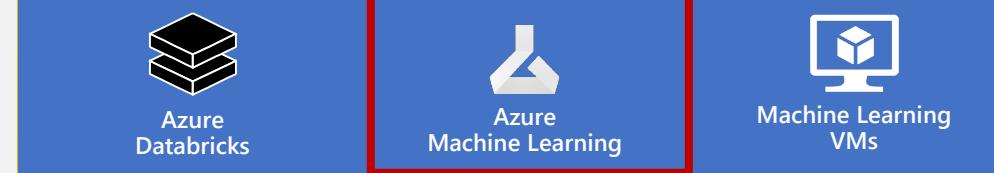
Popular frameworks

To build advanced deep learning solutions



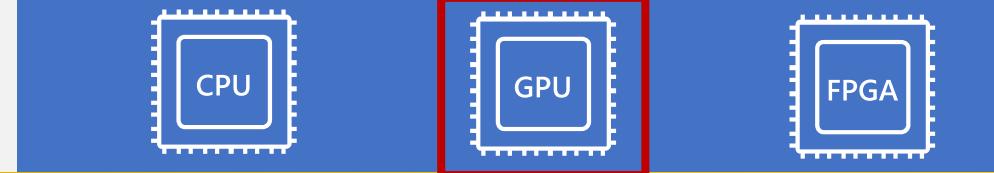
Productive services

To empower data science and development teams



Powerful infrastructure

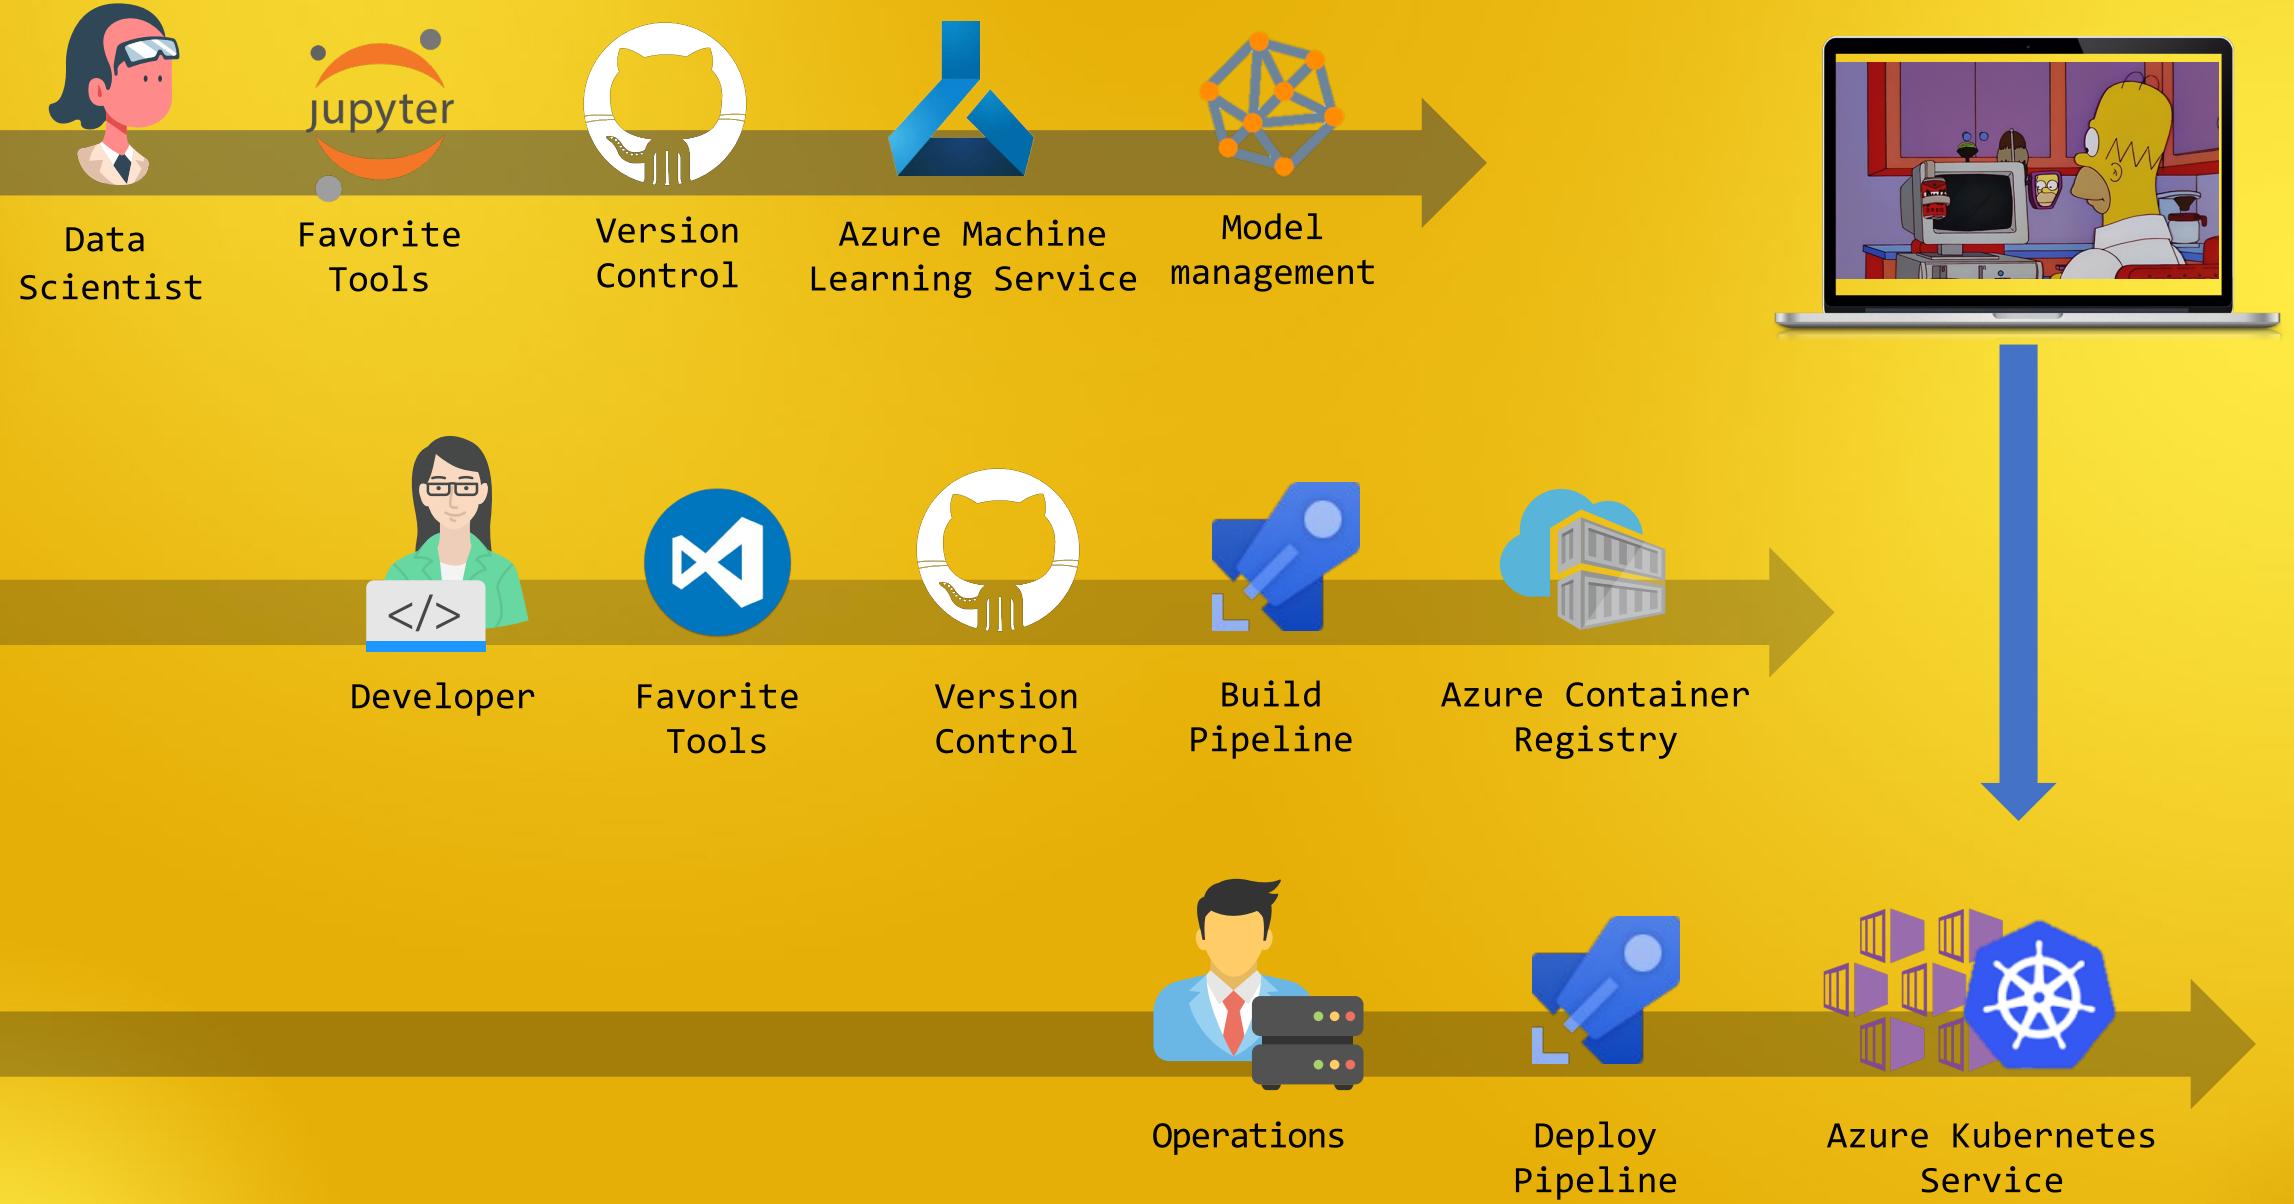
To accelerate deep learning



Flexible deployment

To deploy and manage models on intelligent cloud and edge





What is Azure Machine Learning Service?

Set of Azure
Cloud Services



Python
SDK

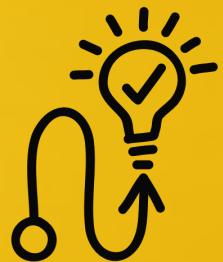
That enables
you to:

- ✓ Prepare Data
- ✓ Build Models
- ✓ Train Models

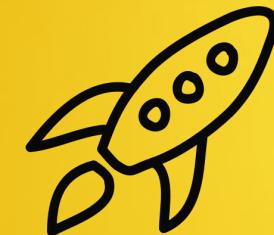
- ✓ Manage Models
- ✓ Track Experiments
- ✓ Deploy Models



Prepare
your environment



Experiment
with your model & data

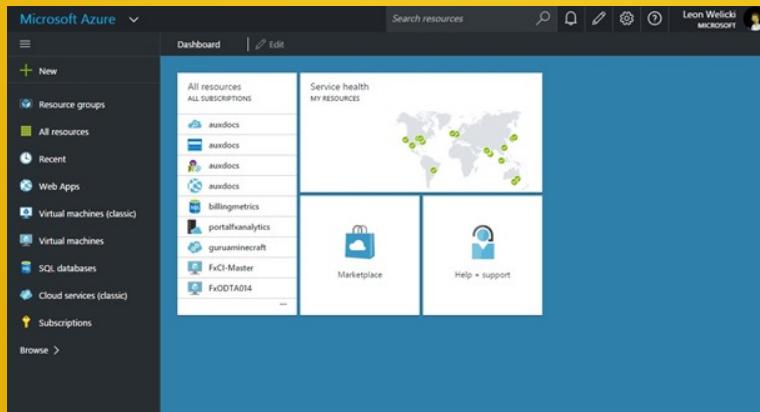


Deploy
Your model into production



Step 1: Prepare your environment

Setup your environment



Azure Portal



VS Code



Python SDK

Create a workspace

```
ws = Workspace.create(  
    name='<NAME>',  
    subscription_id='<SUBSCRIPTION ID>',  
    resource_group='<RESOURCE GROUP>',  
    location='westeurope')
```

```
ws.write_config()
```

```
ws = Workspace.from_config()
```



Create a workspace

Microsoft Azure | Machine Learning

damo-mlworkspace > Home

Welcome!

- Home
- Author
- Automated ML
- Visual Interface
- Notebooks and Files
- Assets
- Datasets
- Experiments
- Models
- Endpoints
- Manage
- Compute
- Datastores
- Notebook VMs

Create New

Automated ML

Automatically train and tune a model using a target metric.

Start Now

Visual Interface

Drag-n-drop interface from prepping data to deploying models.

Start Now

Notebooks

Code with Python SDK and run sample requirements.

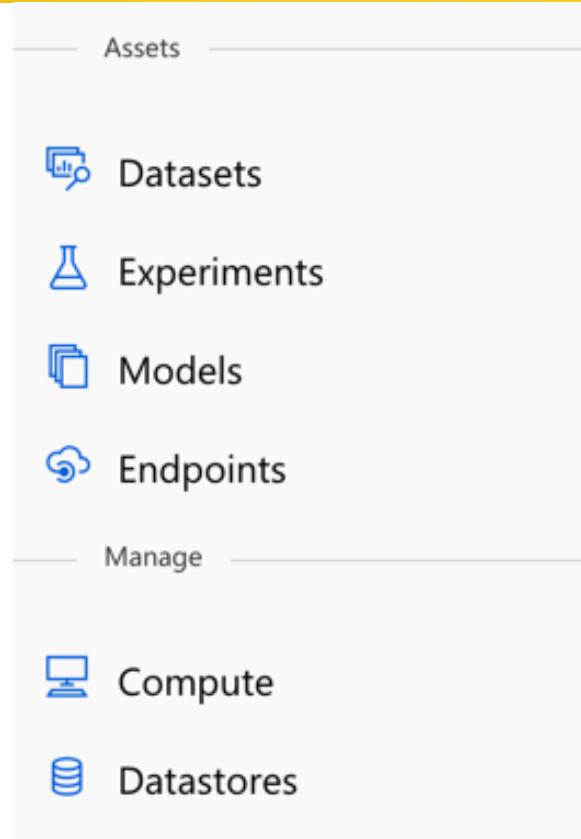
Start Now

My Recent Resources

Run Number	Experiment	Status Updated Time	Status
89	seer	9/12/2019, 12:28:40...	InProg...
85	seer	9/12/2019, 12:07:45...	Failed
81	seer	9/12/2019, 9:53:13 ...	Failed

Name	Type	Provis...
damoseercompute	Machine Learning Compute	Succ...

Azure Machine Learning Service



Datasets – registered, known data sets

Experiments – Training runs

Models – Registered, versioned models

Endpoints:

- Real-time Endpoints – Deployed model endpoints
- Pipeline Endpoints – Training workflows

Compute – Managed compute

Datastores – Connections to data

Create Compute

```
cfg = AmlCompute.provisioning_configuration(  
    vm_size='STANDARD_NC6',  
    min_nodes=1,  
    max_nodes=6)  
  
cc = ComputeTarget.create(ws, '<NAME>', cfg)
```



Create a workspace



Create compute



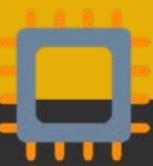
Attach storage

```
ds = ws.get_default_datastore()

ds = Datastore.register_azure_blob_container(
    workspace = ws,
    datastore_name = '<your datastore name>',
    container_name = '<your blob container name>',
    account_name = '<your storage account name>',
    account_key = '<your storage account key>'
)
```



Create a workspace



Create compute



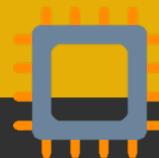
Setup storage



Demo: **Setup your workspace**



Create a workspace



Create compute



Setup storage

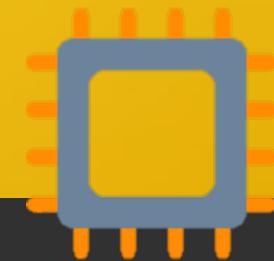


Step 1

Prepare your environment



Create a workspace



Create compute



Setup storage





Step 2: **Experiment with your model & data**

Create an experiment

```
exp = Experiment(workspace=ws, name="<ExperimentName>")
```



Create an
Experiment

Create a training file

```
model=Sequential()
model.add(Conv2D(kernel_size=(3,3),filters=3,input_shape=(128,128,1),activation="relu"))
model.add(Conv2D(kernel_size=(3,3),filters=10,activation="relu",padding="same"))
model.add(MaxPooling2D(pool_size=(2,2),strides=(2,2)))
model.add(Conv2D(kernel_size=(3,3),filters=3,activation="relu"))
model.add(Conv2D(kernel_size=(5,5),filters=5,activation="relu"))
model.add(MaxPooling2D(pool_size=(3,3),strides=(2,2)))
model.add(Conv2D(kernel_size=(2,2),strides=(2,2),filters=10))
model.add(Flatten())
model.add(Dropout(0.3))
model.add(Dense(100,activation="sigmoid"))
model.add(Dense(1,activation="sigmoid"))
model.summary()

model.compile(optimizer="adadelta",loss="binary_crossentropy",metrics=["accuracy"])

history = model.fit(data1, labels, validation_split=0.2, epochs=30, batch_size=20, callbacks=[historyLoss])
```



Create an
Experiment



Create a
training file



Create an Environment

Curated environment

```
curated_env_name = 'AzureML-PyTorch-1.6-GPU'  
pytorch_env = Environment.get(workspace=ws, name=curated_env_name)  
pytorch_env = pytorch_env.clone(new_name='pytorch-1.6-gpu')
```

Custom environment

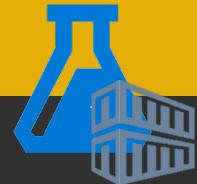
```
env = Environment('pytorch-1.6-gpu')  
cd = CondaDependencies.create(  
    pip_packages=['azureml-dataprep[pandas,fuse]', 'azureml-defaults'],  
    conda_packages=['pytorch', 'torchvision'])  
env.python.conda_dependencies = cd  
env.docker.enabled = True
```



Create an Experiment



Create a training file



Create an ScriptRunConfig



Create an ScriptRunConfig

```
args = ['--data-folder', simpsons_ds.as_named_input('simpsons').as_mount(),
        '--num-epochs', 10]

project_folder = "./trainingscripts"

config = ScriptRunConfig(
    source_directory = project_folder,
    script = 'train.py',
    compute_target=compute_target,
    environment = pytorch_env,
    arguments=args,
)
```



Create an Experiment



Create a training file

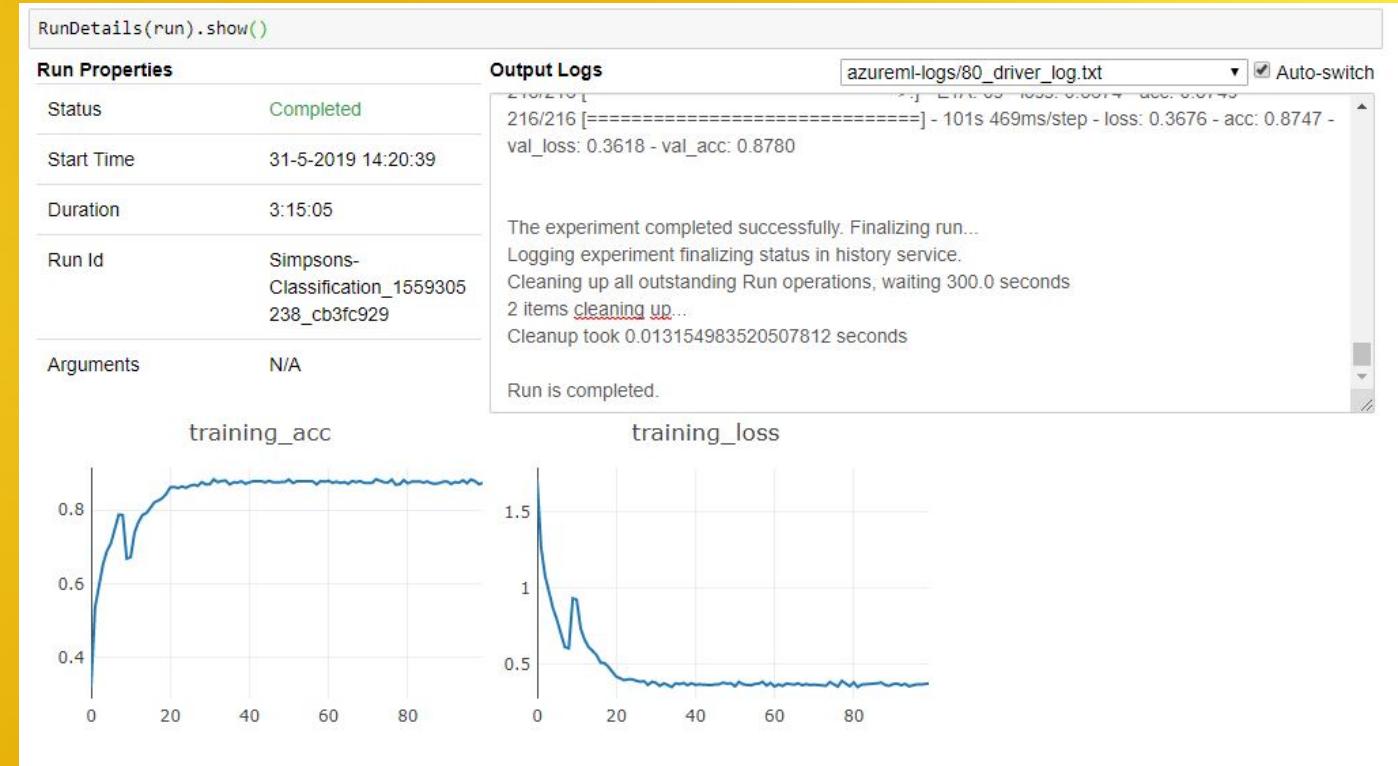


Create an ScriptRunConfig



Submit the experiment to the cluster

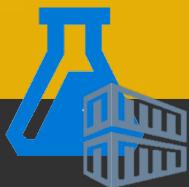
```
run = exp.submit(estimator)  
RunDetails(run).show()
```



Create an Experiment



Create a training file



Create an ScriptRunConfig



Submit to the AI cluster



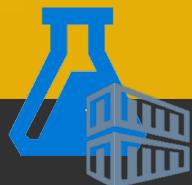
Demo: **Creating and run an experiment**



Create an
Experiment



Create a
training file



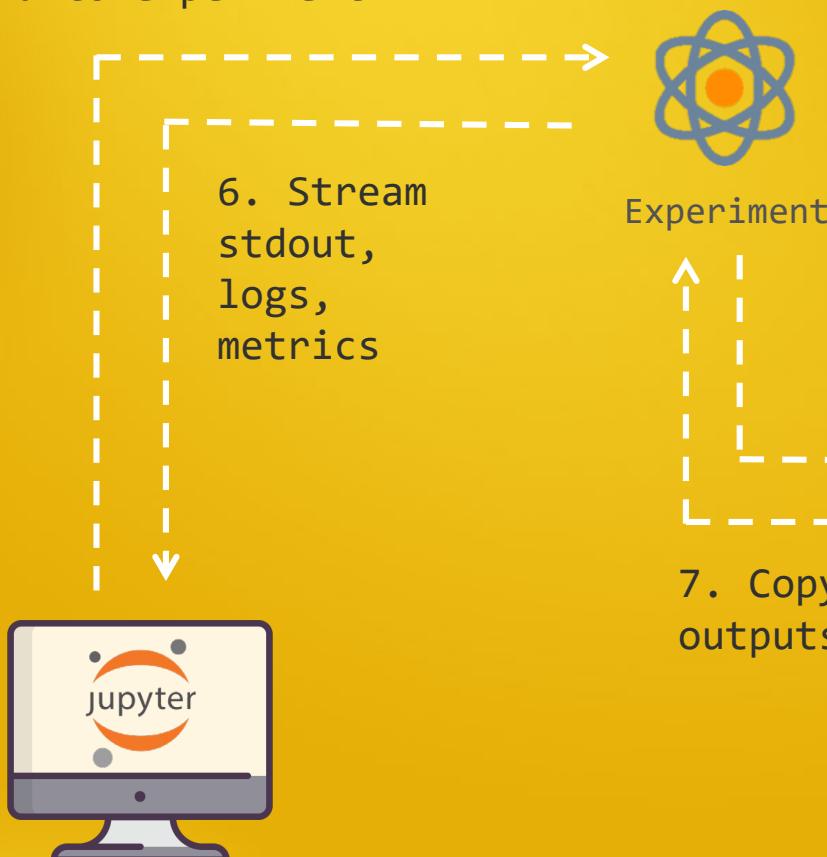
Create an
ScriptRunConfig



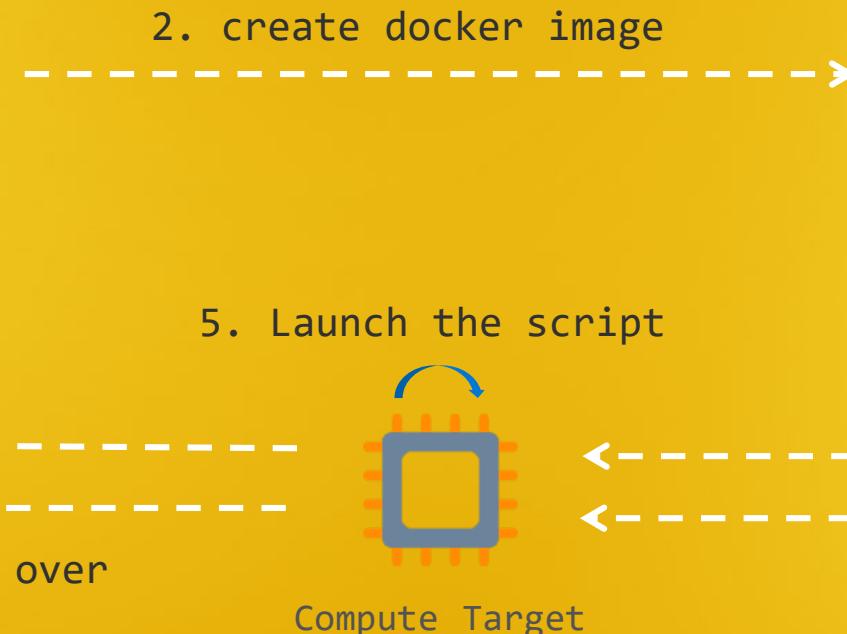
Submit to the
AI cluster



1. Snapshot folder and send to experiment



2. create docker image



Azure Notebook

Experiment

Compute Target

Docker Image

Data store

7. Copy over outputs



6. Stream stdout, logs, metrics

Register the model

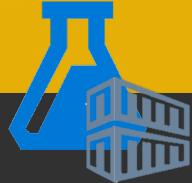
```
model = run.register_model(  
    model_name='SimpsonsAI',  
    model_path='outputs')
```



Create an Experiment



Create a training file



Create an ScriptRunConfig



Submit to the AI cluster



Register the model



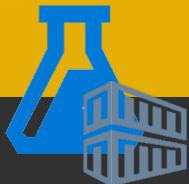
Demo: **Register and test the model**



Create an Experiment



Create a training file



Create an ScriptRunConfig



Submit to the AI cluster



Register the model



Step 2

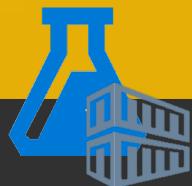
Experiment with your model & data



Create an
Experiment



Create a
training file



Create an
ScriptRunConfig

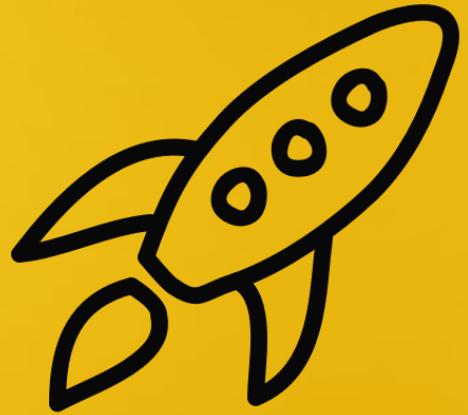


Submit to the
AI cluster



Register the
model





Step 3: Deploy your model into production

AMLS to deploy



The Model



Score.py



Environment file



Docker Image

Score.py

```
%%writefile score.py
from azureml.core.model import Model

def init():
    model_root = Model.get_model_path('MyModel')
    loaded_model = model_from_json(loaded_model_json)
    loaded_model.load_weights(model_file_h5)

def run(raw_data):
    url = json.loads(raw_data)['url']
    image_data = cv2.resize(image_data, (96,96))
    predicted_labels = loaded_model.predict(data1)
    return json.dumps(predicted_labels)
```

Environment File

```
from azureml.core.runconfig import CondaDependencies  
  
cd = CondaDependencies.create()  
cd.add_conda_package('keras==2.2.2')  
cd.add_conda_package('opencv')  
cd.add_tensorflow_conda_package()  
cd.save_to_file(base_directory='./', conda_file_path='myenv.yml')
```

```
1 # Conda environment specification. The dependencies defined in this file will  
2 # be automatically provisioned for runs with userManagedDependencies=False.  
3  
4 # Details about the Conda environment file format:  
5 # https://conda.io/docs/user-guide/tasks/manage-environments.html#create-env-file-manually  
6  
7 name: project_environment  
8 dependencies:  
9     # The python interpreter version.  
10    # Currently Azure ML only supports 3.5.2 and later.  
11    - python=3.6.2  
12  
13    - pip:  
14        - azureml-defaults==1.0.2  
15    - numpy  
16    - keras==2.2.2  
17    - opencv  
18    - tensorflow=1.10.0  
19
```

Inference config

```
inference_config = InferenceConfig(  
    runtime= "python",  
    entry_script="score.py",  
    conda_file="myenv.yml"  
)
```

Deployment using AMLS



Deployment



Azure Container
Instance (ACI)



Azure Kubernetes
Service (AKS)

Deploy to ACI

```
aciconfig = AciWebservice.deploy_configuration(  
    cpu_cores = 1,  
    memory_gb = 2)  
  
service = Model.deploy(workspace=ws,  
    name='simpsons-aci',  
    models=[model],  
    inference_config=inference_config,  
    deployment_config=aciconfig)
```

Deploy to AKS

```
aks_target = AksCompute(ws, "AI-AKS-DEMO")

deployment_config = AksWebservice.deploy_configuration(
    cpu_cores = 1,
    memory_gb = 1)

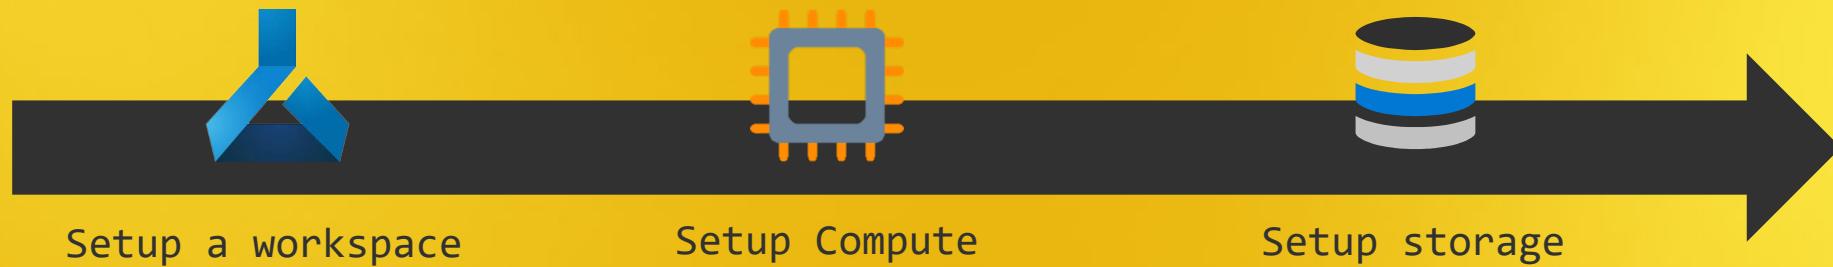
service = Model.deploy(workspace=ws,
                      name="simpsons-ailive",
                      models=[model],
                      inference_config=inference_config,
                      deployment_config=deployment_config,
                      deployment_target=aks_target)

service.wait_for_deployment(show_output = True)
```

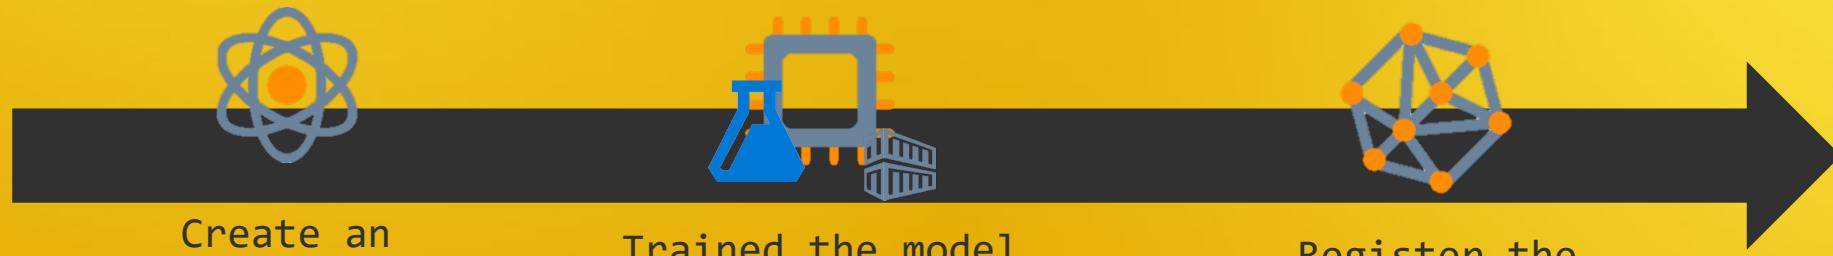
Demo: **Deploy to ACI**

Recap

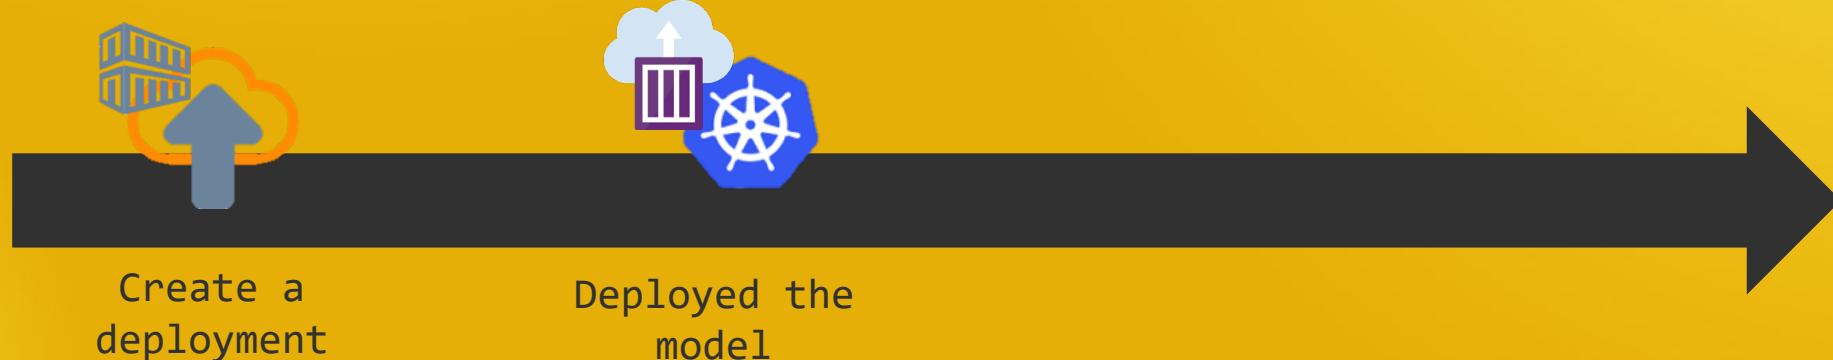
Prepare
your environment



Experiment
with your model & data



Deploy
your model



Code on: **aka.ms/amls-pytorch**

Thank you!

Read more on: aka.ms/mls-pytorch



@hboelman



github.com/hnky



henkboelman.com

